

Cody JUNXIAO WANG

AIGC Researcher | Developer/Engineer Python/PyTorch

✉ jxiao.wang 🌐 github.com/wangjunxiao in linkedin.com/in/junxiao-wang
☎ (+86) 13478902794 ✉ wangjunxiao@live.com 🎓 scholar.google.com
📍 Born Weihai, China, 1991



Several years of experience applying machine learning models *efficiently, faithfully, securely, and privately* in distributed systems. Enjoys designing better problem-solving methods for challenging tasks, and learning new technologies and tools, and contributing to open source and tech communities by sharing knowledge and experience.

Broadly interested in artificial intelligence system with a special focus on generative AI, distributed machine learning, AI security, privacy and interpretability.

📖 Latest Highlight

> Differentially Private Pretraining and Efficient Adaptation for LLMs:

(1) FlashDP: Memory-Efficient and High-Throughput DP-SGD Training for Large Language Models. The NeurIPS Workshop on Adaptive Foundation Models (AFM) 2024, Vancouver, Canada.

TL;DR: Pre-training Llama-13B with 4 A100 GPUs, making the throughput of DP-SGD even reach 90% of NonDP.

(2) ZO-Offloading: Fine-Tuning LLMs with 100 Billion Parameters on a Single GPU. The NeurIPS Workshop on Adaptive Foundation Models (AFM) 2024, Vancouver, Canada.

TL;DR: Fine-tuning extremely large models, such as OPT-175B on a single GPU with just 24GB of memory.

(3) Towards Light Adaptation of Large Language Models For Personal Hardware. The ACM MobiSys Workshop on Edge and Mobile Foundation Models (EdgeFM) 2024, Tokyo, Japan.

TL;DR: Neural operators with efficient hardware inference, high accuracy with no or only a few additional fine-tuning steps.

(4) Autonomous Workflow for Multimodal Fine-Grained Training Assistants Towards Mixed Reality. Findings of the Association for Computational Linguistics (ACL) 2024, Bangkok, Thailand.

TL;DR: Autonomous workflows integrating MLLM agents into LEGO brick assembly in a pilot XR environment.

> Safe Generation for Diffusion Models:

(1) Towards Safe Concept Transfer of Multi-Modal Diffusion via Causal Representation Editing. The Annual Conference on Neural Information Processing Systems (NeurIPS) 2024, Vancouver, Canada.

TL;DR: Efficient and flexible content generation by intervening at diffusion timesteps causally linked to unsafe concepts.

(2) Towards Test-Time Refusals via Concept Negation. The Annual Conference on Neural Information Processing Systems (NeurIPS) 2023, New Orleans, LA, USA.



TL;DR: Flexible concept negation via identification of negative concepts on test-time and purification on the feature space.

📋 Profession & Skills

Programming	Python (PyTorch), Golang, Java, C/C++.
Machine Learning	FL, PEFT, ZO Optimization, DPSGD, Contrastive Learning, Adversarial Learning, Diffusion.
CV/NLP/Multimodal	Transformer, GPT, BERT, ViT, CLIP, CNN, U-Net, GAN, VAE, GNN.
PC/Reviewers	CVPR25, AISTATS25, IJCAI24, ECCV24, AISTATS24, ICCV23, TC, TMC, TIFS.
Misc Tools	Git, LaTeX, Markdown, Linux.

📁 Projects & Experiences

Mar 2024 Feb 2023	Pre-Trained Models, KAUST, Python (PyTorch) <ul style="list-style-type: none">> Implementation of <i>Faithful Interpretation</i> for Concept Bottleneck models.> Implanting triggerable but invisible <i>Trojans</i> onto BERT via random encoding perturbations.> Rebalancing the <i>Slow-Learning Modalities</i> towards the Prototypes.> Implementation of <i>Diffusion Concept Negation</i> with test-time attention refinement.> Publications: KDD23, NeurIPS23, CVPR23, ICLR24.> 🎯 FVLC, Modal-Imbalance-PMR, TrojanAttack. <div>Faithful Interpretation Concept Bottleneck Trojan Modality Rebalance Diffusion</div>
Mar 2023 Mar 2021	Distributed Machine Learning, HONG KONG POLYTECHNIC UNIVERSITY, Python (PyTorch) <ul style="list-style-type: none">> <i>Gradient Inversion</i> protection based on Random Matrix theory.> <i>Knowledge Editing</i> based on TF-IDF and Filter Pruning implemented with NNI.> Implementation of CLIP based <i>FL Framework</i> with Parameter-Efficient Fine-Tuning.> Publications: INFOCOM22, IJCAI22, WWW22, WWW23, TMC, Network.> 🎯 GradDefense, Unlearning, PromptFL. <div>Gradient Random Matrix TF-IDF Filter Pruning NNI CLIP Parameter-Efficient Fine-Tuning</div>

Sep 2019	Network Intelligence, QMUL, Python/C/C++
Sep 2018	<ul style="list-style-type: none"> > Online <i>Traffic Recognition</i> based on 1D-CNN implemented with TensorFlow and Keras. > Implementation of adaptive <i>Sketch Memory Allocation</i> based on Actor-Critic Framework. > Implementation of <i>RFID Integrity Authentication</i> with protocol redesign. > Publications: ICA3PP19, CFI19, SmartIoT19, IoTJ, TMC. >  TrafficCategorization, HBLSketch, RL_MemoryAllocation.
	<div>Traffic Recognition</div> <div>Sketch</div> <div>Memory Allocation</div> <div>Actor-Critic</div> <div>RFID</div>
Dec 2020	Network Systems, DALIAN UNIVERSITY OF TECHNOLOGY, Python/Golang/Java/C/C++
Sep 2016	<ul style="list-style-type: none"> > Implementation of efficient <i>Software Upgrade</i> with State-Isolated Modular Management. > Implementation of fine-grained <i>Control Plane Scheduling</i> with Queue Management. > <i>Data Plane Flow Tracing</i> based on Probabilistic Packet Tagging implemented with OpenFlow. > Publications: SIGCOMM18 Demo, ISJ, TNSM, TNSE, IoTJ, OJCOMS. >  CLICK-UP, FlowTracer, NFVCloud, SDNCloud, SDNDashboard, AgileScheduler.
	<div>Software Upgrade</div> <div>Control Plane</div> <div>Scheduling</div> <div>Queue</div> <div>Data Plane</div> <div>Flow Tracing</div> <div>Packet Tagging</div>

Employment

2023-now	Postdoctoral Fellow, King Abdullah University of Science and Technology (KAUST), Thuwal, Saudi Arabia.
2021-2023	Postdoctoral Fellow, Hong Kong Polytechnic University (PolyU), Hong Kong, China.
2018-2019	Visitor, Queen Mary University of London (QMUL), London, United Kingdom.

Publications

- > Liangyu Wang, Jie Ren, Hang Xu, **Junxiao Wang**, David E. Keyes, Di Wang. ZO-Offloading: Fine-Tuning LLMs with 100 Billion Parameters on a Single GPU. The NeurIPS Workshop on Adaptive Foundation Models (AFM) 2024, Vancouver, Canada.
- > Liangyu Wang, **Junxiao Wang**, Jie Ren, Zihang Xiang, David E. Keyes, Di Wang. FlashDP: Memory-Efficient and High-Throughput DP-SGD Training for Large Language Models. The NeurIPS Workshop on Adaptive Foundation Models (AFM) 2024, Vancouver, Canada.
- > Peiran Dong, Bingjie Wang, Song Guo, **Junxiao Wang**, Jie Zhang, Zicong Hong. Towards Safe Concept Transfer of Multi-Modal Diffusion via Causal Representation Editing. The Annual Conference on Neural Information Processing Systems (NeurIPS) 2024, Vancouver, Canada. (acceptance rate~25.8% [4,043/15,671])
- > Tao Guo, Song Guo, **Junxiao Wang**. Explore and Cure: Unveiling Sample Effectiveness with Context-Aware Federated Prompt Tuning. IEEE Transactions on Mobile Computing (TMC) 2024.
- > Jiahuan Pei, Irene Viola, Haochen Huang, **Junxiao Wang**, Moonisa Ahsan, Fanghua Ye, Jiang Yiming, Yao Sai, Di Wang, Zhumin Chen, Pengjie Ren, Pablo Cesar. Autonomous Workflow for Multimodal Fine-Grained Training Assistants Towards Mixed Reality. Findings of the Association for Computational Linguistics (ACL) 2024, Bangkok, Thailand. (acceptance rate~22.1% [974/4,407])
- > Yizhi Zhou, **Junxiao Wang** (Corresponding Author), Xiangyu Kong, Shan Wu, Xin Xie, Heng Qi. Exploring Amplified Heterogeneity Arising from Heavy-Tailed Distributions in Federated Learning. IEEE Transactions on Mobile Computing (TMC) 2024.
- > Liangyu Wang, **Junxiao Wang**, Di Wang. Towards Light Adaptation of Large Language Models For Personal Hardware. The ACM MobiSys Workshop on Edge and Mobile Foundation Models (EdgeFM) 2024, Tokyo, Japan.
- > Leijie Wu, Song Guo, Yaohong Ding, **Junxiao Wang**, Wenchao Xu, Yufeng Zhan, Anne-Marie Kermarrec. Rethinking Personalized Client Collaboration in Federated Learning. IEEE Transactions on Mobile Computing (TMC) 2024.
- > Songning Lai, Lijie Hu, **Junxiao Wang**, Laure Berti-Equille, Di Wang. Faithful Vision-Language Interpretation via Concept Bottleneck Models. International Conference on Learning Representations (ICLR) 2024, Vienna, Austria. (acceptance rate~31.0% [2,252/7,262])
- > Peiran Dong, Song Guo, **Junxiao Wang** (Corresponding Author), Bingjie Wang, Jiewei Zhang, Ziming Liu. Towards Test-Time Refusals via Concept Negation. The Annual Conference on Neural Information Processing Systems (NeurIPS) 2023, New Orleans, LA, USA. (acceptance rate~26.1% [3,222/12,343])
- > Tao Guo, Song Guo, **Junxiao Wang** (Corresponding Author), Xueyang Tang, Wenchao Xu. PromptFL: Let Federated Participants Cooperatively Learn Prompts Instead of Models - Federated Learning in Age of Foundation Model. IEEE Transactions on Mobile Computing (TMC) 2023.
- > Peiran Dong, Song Guo, **Junxiao Wang** (Corresponding Author). Investigating Trojan Attacks on Pre-trained Language Model-powered Database Middleware. ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD) 2023, Long Beach, CA, USA. (acceptance rate~22.1% [313/1,416])
- > Yunfeng Fan, Wenchao Xu, Haozhao Wang, **Junxiao Wang**, Song Guo. PMR: Prototypical Modal Rebalance for Multimodal Learning. IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) 2023, Vancouver, Canada. (acceptance rate~25.8% [2,360/9,155])
- > Tao Guo, Song Guo, **Junxiao Wang** (Corresponding Author). pFedPrompt: Learning Personalized Prompt for Vision-Language Models in Federated Learning. The ACM Web Conference (WWW) 2023, Austin, Texas, USA. (acceptance rate~19.2% [365/1,900])
- > Leijie Wu, Song Guo, **Junxiao Wang** (Corresponding Author), Zicong Hong, Jie Zhang, Yaohong Ding. Federated Unlearning: Guarantee the Right of Clients to Forget. IEEE Network 2022.
- > Xin Xie, Xiulong Liu, **Junxiao Wang**, Song Guo, Heng Qi, Keqiu Li. Efficient Integrity Authentication Scheme for Large-scale

- RFID Systems. IEEE Transactions on Mobile Computing (TMC) 2022.
- > Rui Zhang, Song Guo, **Junxiao Wang** (Corresponding Author), Xin Xie, Dacheng Tao. A Survey on Gradient Inversion: Attacks, Defenses and Future Directions. International Joint Conference on Artificial Intelligence (IJCAI) 2022, Vienna, Austria. (acceptance rate~14.9% [679/4,535])
 - > **Junxiao Wang**, Song Guo, Xin Xie, Heng Qi. Federated Unlearning via Class-Discriminative Pruning. The ACM Web Conference (WWW) 2022, Online. (acceptance rate~17.7% [323/1,822])
 - > **Junxiao Wang**, Song Guo, Xin Xie, Heng Qi. Protect Privacy from Gradient Leakage Attack in Federated Learning. IEEE International Conference on Computer Communications (INFOCOM) 2022, Online. (acceptance rate~19.8% [224/1,129])
 - > Heng Qi, **Junxiao Wang**, Wenxin Li, Yuxin Wang, Tie Qiu. A Blockchain-driven IIoT Traffic Classification Service for Edge Computing. IEEE Internet of Things Journal (IoTJ) 2021.
 - > **Junxiao Wang**, Heng Qi, Wenxin Li, Keqiu Li, Steve Uhlig, Yuxin Wang. Dynamic SDN Control Plane Request Assignment in NFV Datacenters. IEEE Transactions on Network Science and Engineering (TNSE) 2021.
 - > **Junxiao Wang**, Heng Qi, Keqiu Li, Steve Uhlig. Click-UP: Towards the Software Upgrade of Click based Modular Network Function. IEEE Systems Journal (ISJ) 2020.
 - > Xinping Xu, Wenxin Li, Heng Qi, **Junxiao Wang**, Keqiu Li. Latency-Constrained Cost-Minimized Request Allocation for Geo-distributed Cloud Services. IEEE Open Journal of the Communications Society (OJCOMS) 2020.
 - > **Junxiao Wang**, Heng Qi, Yang He, Wenxin Li, Keqiu Li. FlowTracer: An Effective Flow Trajectory Detection Solution Based on Probabilistic Packet Tagging in SDN-Enabled Networks. IEEE Transactions on Network and Service Management (TNSM) 2019.
 - > Keyan Zhao, **Junxiao Wang**, Heng Qi, Xin Xie, Keqiu Li. HBL-Sketch: A New Three-tier Sketch for Accurate Network Measurement. International Conference on Algorithms and Architectures for Parallel Processing (ICA3PP) 2019, Melbourne, Australia.
 - > Wenrui Zhou, Yuan Cao, Heng Qi, **Junxiao Wang**. An Effective Network Intrusion Detection Framework Based on Learning to Hash. IEEE International Conference on Smart Internet of Things (SmartIoT) 2019, Tianjin, China.
 - > Wanqian Zhang, **Junxiao Wang**, Sheng Chen, Heng Qi, Keqiu Li. A Framework for Resource-aware Online Traffic Classification Using CNN. International Conference on Future Internet Technologies (CFI) 2019, Phuket, Thailand.
 - > **Junxiao Wang**, Heng Qi, Keqiu Li, Xiaobo Zhou. Real-Time Link Fault Detection as a Service for Datacenter Network. Journal of Computer Research and Development 2018.
 - > **Junxiao Wang**, Heng Qi, Keqiu Li, Xiaobo Zhou. PRSFC-IoT: A Performance and Resource Aware Orchestration System of Service Function Chaining for Internet of Things. IEEE Internet of Things Journal (IoTJ) 2018.
 - > **Junxiao Wang**, Yuchen Huang, Heng Qi, Keqiu Li, Steve Uhlig. CLICK-UP: Towards Software Upgrades of Click-driven Stateful Network Element. ACM SIGCOMM Conference 2018 Demo, Budapest, Hungary.

Education

- | | |
|------|---|
| 2020 | PhD in Computer Science, Dalian University of Technology (DUT), Dalian, China. |
| 2017 | MEng in Computer Science, Dalian University of Technology (DUT), Dalian, China. |
| 2014 | BEng in Software Engineering, Dalian Maritime University (DMU), Dalian, China. |